

Guided Lab: Creating an AWS Lambda function

Description

AWS Lambda is a serverless computing service that lets you run code without provisioning or managing servers. Unlike traditional server-based setups like EC2, where you're charged as long as your server is up, AWS Lambda matches costs directly with your actual usage. You pay only when your code is running, which is a big plus for event-based processes that don't need to run 24/7. AWS Lambda automatically scales out, making it a great fit for workloads that have unpredictable or varying demands.

In this lab, you will learn how to create a Lambda function, which is a piece of code that runs in response to events and automatically manages the computing resources required by that code. The lab will guide you through the process of setting up a Lambda function in the AWS Console and testing the function to ensure it works as expected.

Objectives

In this lab, you will:

- Have a basic understanding of Lambda functions
- Understand the cost model of AWS Lambda
- Create and run Lambda functions in the AWS Console using Python
- Understand the general configurations of a Lambda function

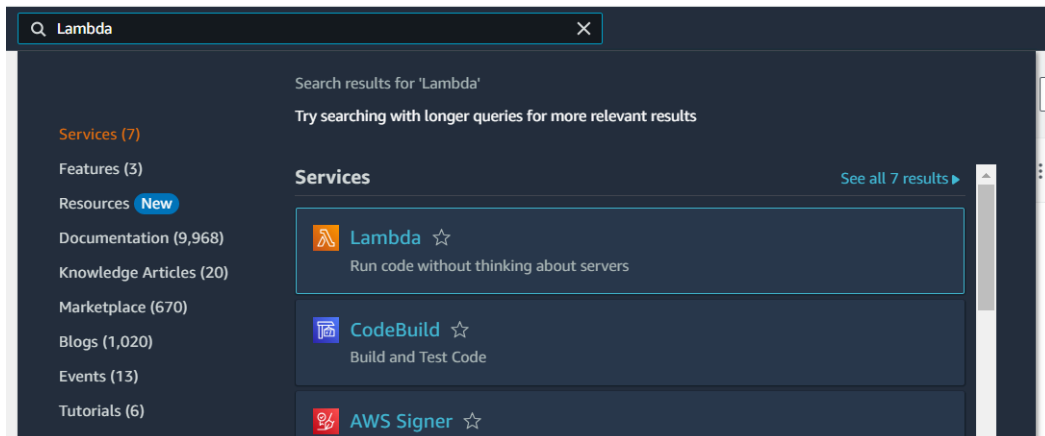
Subscribe to access AWS

PlayCloud Labs


Lab Steps

Create and run Lambda functions in the AWS Console using Python

1. Navigate to the search bar, type “Lambda”, and click to open the Lambda.

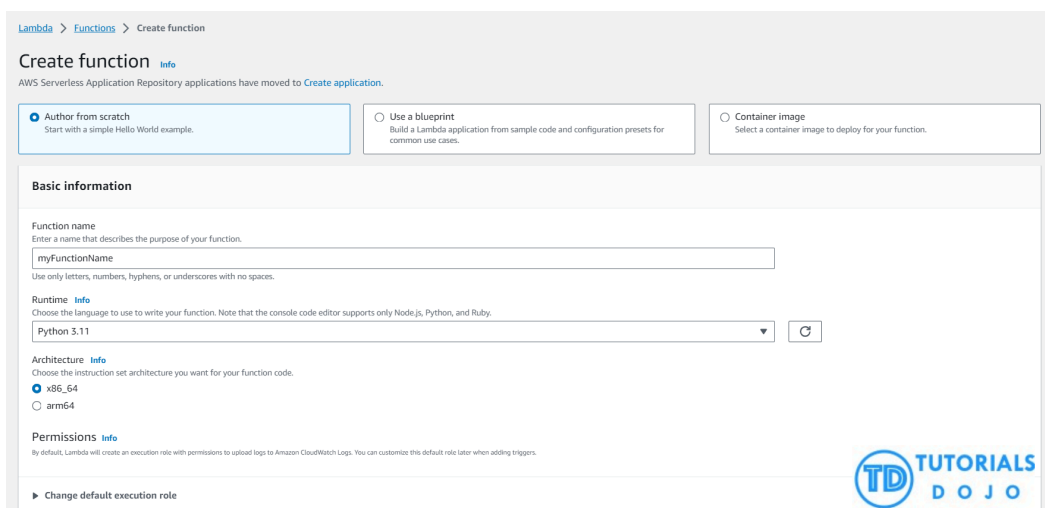


2. Click on the “Create function” button. You will be prompted to configure your function.



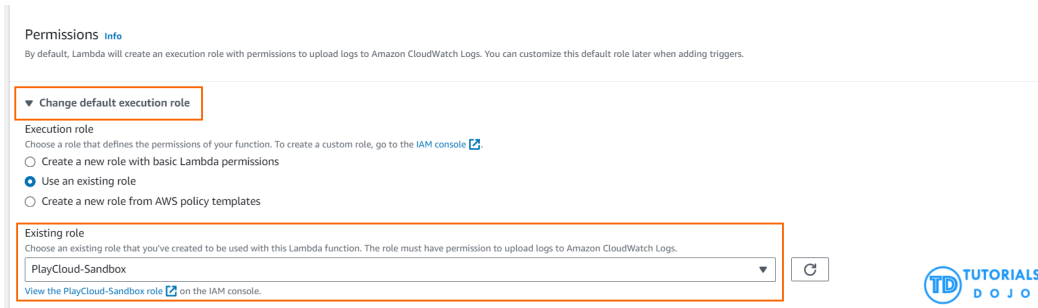
3. Configure your function:

- **Function name:** Give your function a unique name.
- **Runtime:** Select “Python 3.11” from the dropdown menu.



4.To change the default execution role, please follow these steps:

- Go to the “Permissions” tab.
- Navigate to the section where you can edit the execution role.
- Modify the default execution role to “Use an existing role”.
- Choose the existing role named “PlayCloud-Sandbox”.

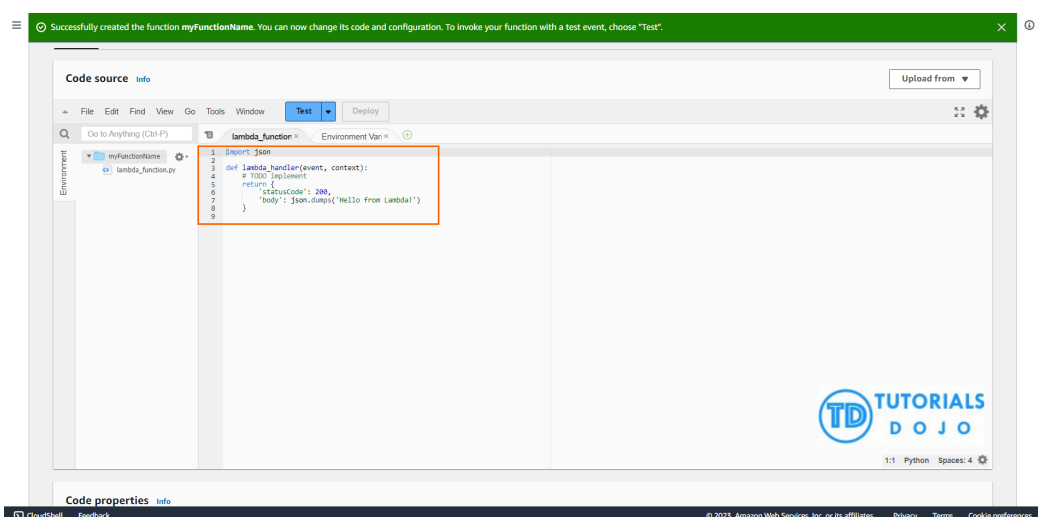


An execution role grants the Lambda function permissions to access AWS services and resources. For example, if you want your Lambda function to make API calls to access an S3 bucket, its execution role would need to include policies allowing it to read and write to that specific bucket. In this lab, an execution role with the necessary permissions has already been created for you.

5. Click “Create Function”

6. In the Function code section, you will see an inline code editor where you can write your Python code. AWS provides a simple “Hello from Lambda!” example to get you started.

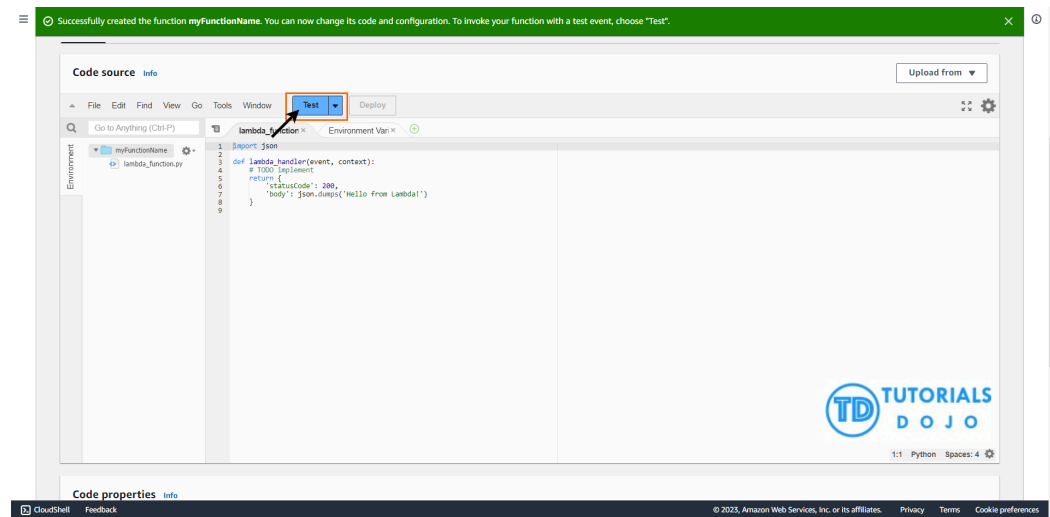
Note: We are using the **old console editor** for this lab. You can switch to the **new editor** as you desire; the process remains the same, but the interface may look slightly different.



7. If you have changes to the code, click on the “Deploy” button at the top to save your function.

8. Test your function:

- Click on the “Test” button at the top of the page.



- In the Configure test event dialog, give your test event a name, and provide a JSON-formatted input for your function. Then, click on “Save” to close the dialog and create your test event.

Configure test event

A test event is a JSON object that mocks the structure of requests emitted by AWS services to invoke a Lambda function. Use it to see the function's invocation result.

To invoke your function without saving an event, configure the JSON event, then choose Test.

Test event action

☒ Create new event
 ☐ Edit saved event

Event name

1

Maximum of 25 characters consisting of letters, numbers, dots, hyphens and underscores.

Event sharing settings

☒ Private

This event is only available in the Lambda console and to the event creator. You can configure a total of 10. [Learn more](#)

☐ Shareable

This event is available to IAM users within the same account who have permissions to access and use shareable events. [Learn more](#)

Template - optional

Event JSON

Format JSON

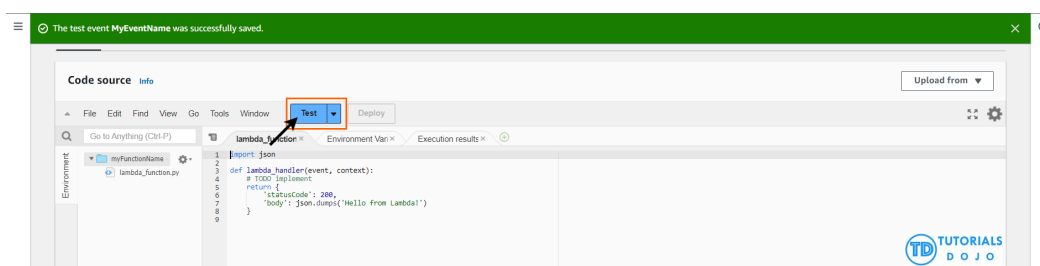
```

1 {
2   "key1": "value1",
3   "key2": "value2",
4   "key3": "value3"
5 }

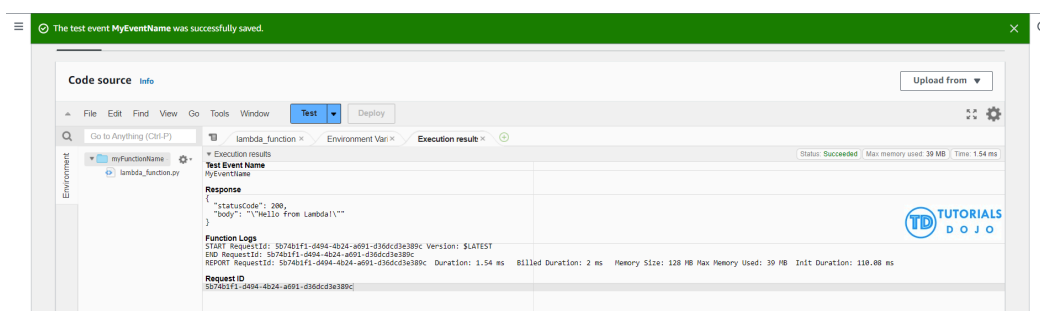
```

2

- Back in the main console, click on "Test" again to run your function with your test event.




9. After running your function, you will see the execution results at the top of the page. This includes Test Event Name, Response, Function Logs, Duration of execution, and Request ID.



General configuration

1. Go to the "Configuration" tab.



2. **Ephemeral Storage** – This refers to the temporary storage space that a Lambda function has for reading or writing data. By default, AWS Lambda allocates 512 MB for a function's /tmp directory. You can increase or decrease this amount using the "Ephemeral storage (MB)" setting¹. To configure the size of a function's /tmp directory, you can set a whole number value between 512 MB and 10,240 MB, in 1-MB increments. This feature allows you to control the amount of ephemeral storage a function gets, enabling you to use AWS Lambda for ETL jobs, ML inference, or other data-intensive workloads.

3. **SnapStart** – is a performance optimization feature developed by AWS that can significantly improve the startup time for applications. Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code. The largest contributor to startup latency (often referred to as cold start time) is the time that Lambda spends initializing the function, which includes loading the function's code, starting the runtime, and initializing the function code. With SnapStart, Lambda initializes your function when you publish a function version. Lambda takes a Firecracker microVM snapshot of the memory and disk state of the initialized execution environment, encrypts the snapshot, and caches it for low-latency access. When you invoke the function version for the first time, and as the invocations scale up, Lambda resumes new execution environments from the cached snapshot instead of initializing them from scratch, improving startup latency.

4. **Memory** – refers to the amount of memory available to your Lambda function at runtime.

You can increase or decrease the memory and CPU power allocated to your function using the “Memory (MB)” setting. To configure the memory for your function, you can set a value between 128 MB and 10,240 MB in 1-MB increments. Lambda allocates CPU power in proportion to the amount of memory configured. Thus, increasing the memory also increases the CPU power, thereby increasing the overall computational power available.

To configure the memory for your function, you can follow these steps:

- Assume that we’re executing this particular Lambda Function. This code finds all prime numbers up to a certain limit.

```
import json

def lambda_handler(event, context):
    upper_limit = int(event.get('upper_limit', 1000000))
    primes = find_prime(upper_limit)
    return {
        'statusCode': 200,
        'body': json.dumps({'message': f'Found {len(primes)} primes up to {upper_limit}.'})
    }

def find_prime(n):
    sieve = [True] * (n + 1)
    sieve[0:2] = [False, False]
    for current in range(2, int(n**0.5) + 1):
        if sieve[current]:
            sieve[current*current::current] = [False] * len(range(current*current, n+1, current))
    return [i for i, v in enumerate(sieve) if v]
```

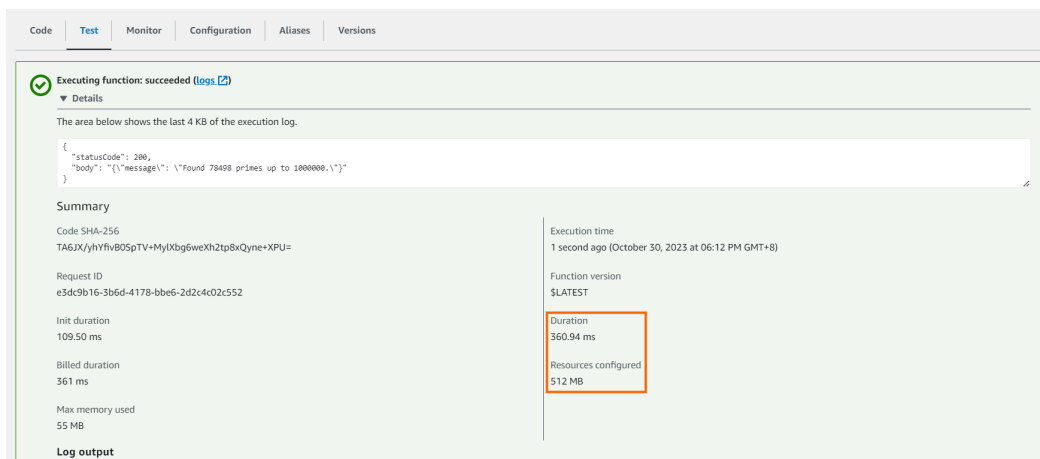
- To save the changes made in the code, click on the “Deploy” button.
- To test your code, head over to the “Test” tab and click on the “Test” button.

- Take note of the result.

- On the function configuration page, on the General configuration pane, choose Edit.

- In the “Memory (MB)” field, increase the memory from 128 MB to 512 MB.

- Click on the “Save” button to apply your changes.
- Head over to the “Test” tab again and click on the “Test” button.
- You will notice that the function executed faster than the first time we ran it, as indicated by the shorter duration.



5. Timeout – refers to the maximum amount of time that your Lambda function can run before it is stopped by the Lambda services.

You can set the timeout value between 1 and 900 seconds (15 minutes). The default timeout value in the Lambda console is 3 seconds. This timeout value acts as a safety buffer that prevents functions that never exit from running indefinitely.

To set the timeout for an AWS Lambda function, you can follow these steps:

- Copy and paste the following code into your Lambda function. This code simulates a function performing actual processing by pausing the execution for 5 seconds. Since the default timeout is 3 seconds, the function will return a timeout error.


```
import json
import time

def lambda_handler(event, context):
    # TODO implement
    time.sleep(5);
    return {
        'statusCode': 200,
        'body': json.dumps('Hello from Lambda!')
    }
```

- To save the changes made in the code, click on the “Deploy” button.
- Upon execution, the function will display an error message similar to the screenshot below:

```
Response
{
  "errorMessage": "2023-10-30T09:45:01.061Z 40a71a9d-76e7-4b9b-a436-f9708ea8d0fa Task timed out after 3.01 seconds"
}
```

- On the General configuration pane, choose Edit.



- For Timeout, set a value from 3 seconds to 10 seconds.

Timeout

min sec

- Click on the "Save" button to apply your changes.
- Go to the "Code" tab and execute the function again.
- Your function should be able to run now after configuring the timeout for your AWS Lambda function.

It's worth noting that setting a timeout value close to the average duration of a function can increase the risk of unexpected timeouts. This is because the duration of a function can vary depending on factors such as the amount of data to be processed, the latency of the services it interacts with, and the amount of processing required.